

WHAT HIDES MUSICAL GEOMETRY IN AUDIO REPRESENTATIONS?

Junyoung Koh

School of Music, Theatre and Dance, University of Michigan
Department of Artificial Intelligence, Yonsei University

ABSTRACT

Learned representations are often evaluated by how well their attributes can be decoded. However, linear decodability does not imply that a representation’s ambient geometry reflects the relationships among those attributes. Music makes this distinction directly testable because relationships between musical classes, such as keys and pitches, are explicitly defined. We examine this gap through linear decoding, class-centroid geometry, and linear-readout geometry. A log-CQT decodes key better than MERT-v1-330M (.492 vs. .456), yet shows no circle-of-fifths structure in its class centroids while the same structure is strongly expressed in its readout weights. We further find two distinct sources of this mismatch. First, increasing the number of examples per centroid changes fifths geometry by a factor of 13.4, compared with only 1.7 for a thousandfold increase in pooling duration. Second, averaging fails to recover octave equivalence, whereas a rank-4 projection fitted only on GiantSteps keys exposes octave structure in individual NSynth recordings using less than 0.3% of the variance. These results show that decodability, ambient geometry, and readout geometry are distinct properties of a representation, and that musical relations can differ substantially in where and how strongly they appear.

Index Terms— music representation learning, probing, representation geometry, key detection, octave equivalence

1. INTRODUCTION

Self-supervised music audio encoders support linear probes [1] for key, pitch, genre, and instrument [2, 3], and this is usually taken to mean that these attributes are *organized* in the representation. Music provides a useful setting for testing this interpretation because music theory specifies not only which classes exist, but also how they are related. Keys a perfect fifth apart share six of seven scale degrees and are adjacent on the circle of fifths, while pitches an octave apart belong to the same pitch class.

We compare what can be decoded from a representation, how musical classes are arranged by distances in the original

representation space, and how those classes are arranged in the weights of a linear classifier. Across these measurements, linear accessibility does not necessarily correspond to structure in the original representation geometry. Our results trace this mismatch to two distinct sources. For clip-level key geometry, the dominant factor is the number of examples used to estimate each class centroid rather than the duration of each example. For note-level octave equivalence, increasing the number of examples does not recover the structure, whereas a learned projection exposes it under a different metric. Together, these results show that the absence of structure in the original geometry does not by itself imply the absence of that structure from the representation.

2. RELATED WORK

Linear probes are the standard instrument for asking what a representation encodes [1], and music audio encoders are routinely evaluated this way, with MERT [2] and MARBLE [3] reporting probe accuracy for key, pitch, genre and instrument. Hewitt and Liang [4] show that a probe’s success has to be read against a control task that removes the structure of interest, and we adopt that design for representation geometry rather than for probe complexity.

Music cognition has established well-defined relational structures for both key and pitch. Perceived key relationships follow the circle of fifths and Krumhansl–Kessler profiles [5, 6], while pitch perception separates pitch height from chroma [7]. Chroma features [8, 9] encode octave equivalence by construction and therefore provide a positive control for these relationships. Tonality-aware objectives such as STONE [10] explicitly impose key equivariance during training. Pilataki et al. [11] report a related observation in a JEPA-based transcription model, where NSynth embeddings show strong separation by timbre and pitch height but weak separation by pitch class.

3. METHOD

Representation decomposition. We write an item of class y as

$$z = s_y + n_y + \varepsilon, \quad \mathbb{E}[\varepsilon \mid y] = 0, \quad (1)$$

Code, pre-registrations and all run artefacts: <https://github.com/Junst/topo-music>. Every threshold and kill criterion reported here was registered before the corresponding run.




Fig. 1. Clip-level key geometry is set by centroid sample size, not by temporal context. (a) and (b) share a vertical axis, with error bars of one standard deviation over 5 draws in (a) and 10 in (b). MuQ L6 sits at the same fraction of depth as MERT L12. (c) The same rise seen as geometry, for the 24 MERT L12 key centroids ordered round the circle of fifths with majors above the white rule. Each panel is rank-transformed over its own entries and each sampled one is the median of ten draws.

where s_y represents the class-dependent signal whose relations encode the structure of interest, such as key relationships or octave equivalence, n_y captures other variation that depends on the class, and ε captures variation whose conditional mean is zero within each class. Averaging examples from the same class reduces ε but not n_y , whereas a change of metric can alter how s_y and n_y contribute to the measured geometry.

Representations. MERT-v1-330M [2] layers 4 to 24, MuQ-large [12] layers 2 to 12, which matches MERT in size but is trained to predict Mel-RVQ tokens and is half as deep, so its layers 2, 6, 8 and 12 sit at the same fractions of depth as MERT’s 4, 12, 16 and 24, EnCodec 32 kHz [13], as used by MusicGen [14], and a log-CQT [8, 9] as a spectral baseline. A CQT chromagram [9], which folds frequencies across octaves into 12 pitch classes, provides a positive control for octave-invariant structure. Frame-level representations are averaged over time to obtain one vector per audio clip.

Key-centroid geometry. On GiantSteps [15], with 7035 clips covering all 24 keys, we average the vectors of tracks sharing a key into 24 class centroids and rank-correlate their pairwise Euclidean distances against three references: circle-of-fifths distance, the Krumhansl–Kessler key-profile distance, and semitone distance. Significance uses a 2000-draw permutation null over key labels, and a partial Spearman additionally removes mel-spectral centroid distance. Semitone distance is included as a control to distinguish circle-of-fifths structure from proximity in chromatic pitch space.

Readout geometry. The same probe, a multinomial logistic regression under 5-fold cross-validation, yields 24 class weight vectors, and we apply the identical measurement to their pairwise cosine distances.

Octave equivalence. On NSynth [16] we use real recordings only, since pitch shifting by signal processing would build the answer into the stimulus. For each instrument we fix an anchor set of pitches supporting every transposition up to 25 semitones, held constant across transpositions, and measure the mean cosine similarity between a note and its transposition by k semitones. Three statistics are reported together, as registered in advance: a local octave contrast against the mean of 11 and 13 semitones; a strict contrast against the largest similarity from 8 to 11; and the coefficient on a twelve-semitone periodic term in a fitted curve. Confidence intervals are bootstrapped over *instruments*. We require the three to agree in sign, since a positive periodic coefficient with no local peak at 12 semitones is a curve-fitting artefact.

Subspaces. Each subspace is an orthonormal basis of the top d shrinkage-LDA discriminant directions for a stated target, with d at most 11, the rank of any linear map to a twelve-class target. We report the fraction of total variance it holds and, as a control, matched random orthonormal subspaces, since projection changes the metric whatever the directions are.

4. RESULTS

4.1. Information, metric and readout dissociate

Table 1 measures all three on the same clips and keys, and the control block sets the scale. An analytic circle of fifths scores close to 1 in every column, so a null elsewhere is a null of the representation and not of the method, while random features show nothing anywhere. Orthogonal per-key codes are highly decodable with tonal relations absent by construction. They

Table 1. Key information, ambient geometry and readout geometry come apart, on the same 7035 GiantSteps clips and 24 keys. acc is 24-way key accuracy, chance .042; ρ_5^{cent} is the circle-of-fifths correlation among the 24 key centroids and ρ_5^W the same among the probe’s class weight vectors; ρ_{chr} is semi-tone adjacency, a discriminant a generic tendency for nearby keys to look alike would also raise. Partialling out mel-spectral centroid distance changes no sign. Orthogonal key codes decode at .651 with classes equidistant by construction and still show nothing, so ρ_5^W is not an artefact of an accurate classifier. In each of the first three columns blue marks the highest value among the real representations and red the lowest; the control block and ρ_{chr} are excluded.

representation	acc	ρ_5^{cent}	ρ_5^W	ρ_{chr}
<i>Controls</i>				
analytic fifths	0.274	+0.977	+0.986	−0.085
random features	0.055	+0.009	+0.033	−0.001
orthogonal key codes	0.651	−0.008	−0.005	−0.001
<i>Spectral front ends</i>				
log-CQT	0.492	−0.035	+0.560	+0.174
log-CQT, 84 bin, per-frame norm	0.486	−0.117	+0.671	+0.491
octave-summed, 12 bin, dB	0.407	−0.106	+0.305	+0.391
PQ-STFT	0.479	+0.550	+0.172	+0.112
chromagram	0.451	+0.475	+0.692	+0.140
<i>Neural codec</i>				
EnCodec 32 kHz	0.404	+0.177	+0.648	+0.041
<i>MERT-v1-330M</i>				
layer 4	0.451	+0.445	+0.412	−0.025
layer 12	0.456	+0.468	+0.428	+0.024
layer 16	0.440	+0.395	+0.409	+0.047
layer 24	0.418	+0.294	+0.484	+0.049
<i>MuQ-large</i>				
layer 2	0.484	+0.533	+0.103	+0.021
layer 6	0.457	+0.530	+0.316	+0.016
layer 8	0.434	+0.545	+0.364	+0.002
layer 12	0.420	+0.372	+0.464	+0.056

decode key at .651, exceeding every real representation, and yet show no fifths geometry in either their centroids, at -0.008 , or their readout weights, at -0.005 . This follows the logic of control tasks [4] without relying on matched accuracy, and this rules out an alternative explanation: that any sufficiently accurate 24-way classifier acquires fifths-structured weights because neighboring keys are intrinsically confusable.

The dissociation that follows is unambiguous, in that the log-CQT has the highest key accuracy of any real arm and no fifths structure in its centroids, at -0.035 with p of .52, while its probe weights carry that structure strongly, at .560. Tonal relations are therefore linearly accessible in a log-CQT without being expressed in its ambient geometry.

For MERT the main change is in ambient geometry rather than linear accessibility, with readout geometry between 0.41 and 0.48, below the log-CQT, EnCodec and the chromagram, and centroid geometry far higher. Across depth tonal relational geometry weakens, from 0.445 at layer 4 to 0.294 at layer 24, while major and minor become more separable. We report this




Fig. 2. Two views of the same rank-4 projection, fitted only on GiantSteps track tonics and applied unchanged to individual NSynth recordings. (a) Ambient similarity decays smoothly with transposition distance, while the projected curve reaches maxima at exactly twelve and twenty-four semitones. (b) The strict contrast that curve summarises, with the matched random subspace as a control. In MERT ambient and random are both near zero; in MuQ they are equal and above it, so there the learned projection must be read against the random one. Error bars are instrument-bootstrap 95% intervals.

as a correlation rather than a mechanism, since key accuracy does not improve along the way and the mode effect is not significant in the chromagram control.

MuQ supplies the mirror image. Its layer 2 decodes key at .484, within a point of the log-CQT, carries the strongest centroid fifths geometry of any learned arm at 0.533, and yet shows none in its readout weights, at 0.103 with p of .078. Both one-sided cases therefore occur at matched accuracy, and across MuQ’s depth the two exchange. PQ-STFT places the same profile on the other side of the learned and handcrafted divide, with the strongest centroid fifths geometry of any arm at 0.550 against 0.172 in its readout, though its chromatic correlation is also positive, at 0.112 with p of .048, so that centroid effect is not exclusively fifths-specific. Where relational structure appears is therefore not settled by whether a representation was learned.

4.2. Sample-limited masking

The clip-level result invites a temporal reading, in which keys are properties of seconds of music, but that reading does not survive the measurement. Fig. 1(a) sweeps the pooling window

from one frame to the full 20 s clip, drawing one window per clip so the number of clips behind each centroid is held constant. A single 13 ms MERT frame already yields 0.276, which is 59% of the 20 s value, MuQ starts at 66% of its own, and the chromagram is flat to within 3%, with no knee anywhere in the sweep. EnCodec is the exception, rising sixfold, as a codec latent carrying little beyond its own frame would.

The clip-level measurement averages twice, and only one of those was being varied. Fig. 1(b) holds the window at the full clip and sweeps clips per centroid, giving 0.035 at one clip and 0.468 at all of them for MERT L12, a factor of 13.4 against 1.7 for a thousandfold longer window. In the single-example centroid regime, fifths geometry is indistinguishable from zero for every arm, and Fig. 1(c) shows the centroid distance matrices that correlation summarises.

This is Eq. (1) with ε dominant, because what differs between two GiantSteps tracks in the same key, namely production, instrumentation and subgenre, is unrelated to the key and cancels in the mean. Centroid sample size is therefore the dominant source of the apparent clip-level effect, and a longer window accounts for a small fraction of it.

4.3. Projection exposes octave-equivalent structure

Ambient octave equivalence varies across encoders. Under the sign-consistency criterion registered in advance, MuQ satisfies it at every tested layer, though weakly, its strict contrast falling from 0.037 at layer 2 to 0.022 at layer 12 against 0.698 for the chromagram. MERT and the codecs fail it at every depth. Averaging NSynth recordings into pitch centroids recovers no reliable effect in MERT at any sample size, so its near-zero ambient geometry is not sample-limited centroid estimation, which is the behaviour Eq. (1) attributes to n_y .

A rank-4 LDA projection fitted only to GiantSteps tonic labels, with no notes and no octave information, substantially increases octave equivalence on individual recordings. The strict contrast rises from -0.001 to 0.304 for MERT layer 4 and from 0.035 to 0.346 for MuQ layer 2, with similar effects at later layers, and is largest at the shallowest layer tested in both encoders. Fig. 2(a) shows the curve those contrasts are drawn from, where the projected similarities reach maxima at exactly twelve and twenty-four semitones. For MuQ the matched random rank-4 subspace retains about the ambient effect, at 0.039 for layer 2, and stays far below the learned projection, while for MERT it lies between -0.019 and 0.007 . They account for 0.08 to 0.29% of NSynth variance in MuQ and 0.11 to 0.13% in MERT.

Key supervision does not guarantee pitch-height invariance, since GiantSteps key labels may correlate with subgenre, tempo or production, though the transfer rules out an explanation resting only on GiantSteps-specific nuisance correlations. Nor is it mechanistic: whether the projection suppresses class-coupled variation or amplifies the tonal component is not iden-

tifiable here. Where the absence has been read as a property of the pretraining objective [11], the transfer strengthens the relation in a fixed representation without retraining.

A registered prediction that failed. We predicted that removing the recovered subspace would at least halve clip-level fifths geometry. Against a baseline recomputed on the same held-out clips it barely moves, from 0.127 to 0.119 for MERT layer 12 and from 0.289 to 0.281 for MuQ layer 2, and only the twelve-dimensional chromagram collapses, which is a dimensionality effect. The recovered subspace is therefore sufficient to expose the relation but does not exclusively contain it. Establishing how widely it is distributed would need repeated recovery after sequential deflation.

5. DISCUSSION

Information availability, metric visibility and readout accessibility are separable, and separating them changes what a probe result licenses: accuracy says the structure is linearly recoverable, and nothing about whether it governs the distances a similarity or retrieval system consumes.

The two masking modes suggest different remedies. Sample-limited masking is a property of the evaluation, and the fix is more examples per class rather than longer audio, so the number of items behind each centroid should be reported as standard, since the same encoder scores 0.035 or 0.468 depending only on that. Metric masking instead requires changing how distances are measured, and the MuQ control separates that reweighting from dimensionality reduction: random rank-4 subspaces preserve about the ambient octave effect, while tonic-discriminative directions amplify it.

Nothing in Eq. (1) is specific to music, since semantic relations in vision may be linearly available while the geometry is dominated by pose, and phonetic structure in speech may be masked by speaker and channel.

Limitations. GiantSteps is entirely electronic dance music, where key correlates with subgenre, and the cross-dataset transfer constrains that confound without removing it. NSynth is monophonic and drawn from isolated instruments, so its nuisance structure is unusually clean. Our subspaces are linear by design, and a nonlinear transform might expose more at the cost of interpretability.

6. CONCLUSION

Musical structure does not occupy a fixed geometric form across representations. Decodability, ambient geometry and readout geometry vary independently: the log-CQT decodes key well with no fifths structure in its centroids and strong structure in its readout, while PQ-STFT and MuQ show the reverse at matched accuracy. Averaging separates sample-limited estimates from stable metric effects, and a projection fitted on another dataset strengthens octave equivalence averaging never reaches.

7. REFERENCES

- [1] Guillaume Alain and Yoshua Bengio, “Understanding intermediate layers using linear classifier probes,” in *Proc. ICLR Workshop Track*, 2017.
- [2] Yizhi Li, Ruibin Yuan, Ge Zhang, Yinghao Ma, Xingran Chen, Hanzhi Yin, et al., “MERT: Acoustic music understanding model with large-scale self-supervised training,” in *Proc. ICLR*, 2024.
- [3] Ruibin Yuan, Yinghao Ma, Yizhi Li, Ge Zhang, Xingran Chen, Hanzhi Yin, et al., “MARBLE: Music audio representation benchmark for universal evaluation,” in *Proc. NeurIPS Datasets and Benchmarks Track*, 2023.
- [4] John Hewitt and Percy Liang, “Designing and interpreting probes with control tasks,” in *Proc. EMNLP*, 2019, pp. 2733–2743.
- [5] Carol L. Krumhansl and Edward J. Kessler, “Tracing the dynamic changes in perceived tonal organization in a spatial representation of musical keys,” *Psychol. Rev.*, vol. 89, no. 4, pp. 334–368, 1982.
- [6] Carol L. Krumhansl, *Cognitive Foundations of Musical Pitch*, Oxford University Press, New York, 1990.
- [7] Roger N. Shepard, “Geometrical approximations to the structure of musical pitch,” *Psychol. Rev.*, vol. 89, no. 4, pp. 305–333, 1982.
- [8] Judith C. Brown, “Calculation of a constant Q spectral transform,” *J. Acoust. Soc. Amer.*, vol. 89, no. 1, pp. 425–434, 1991.
- [9] Brian McFee, Colin Raffel, Dawen Liang, Daniel P.W. Ellis, Matt McVicar, et al., “librosa: Audio and music signal analysis in Python,” in *Proc. 14th Python in Science Conf.*, 2015.
- [10] Yuexuan Kong, Vincent Lostanlen, Gabriel Meseguer-Brocal, Stella Wong, Mathieu Lagrange, and Romain Hennequin, “STONE: Self-supervised tonality estimator,” in *Proc. ISMIR*, 2024.
- [11] Mary Pilataki, Matthias Mauch, and Simon Dixon, “Self-supervised representation learning with a JEPA framework for multi-instrument music transcription,” in *Proc. IEEE WASPAA*, 2025.
- [12] Haina Zhu, Yizhi Zhou, Hangting Chen, Jianwei Yu, Ziyang Ma, Rongzhi Gu, et al., “MuQ: Self-supervised music representation learning with mel residual vector quantization,” *IEEE Trans. Audio, Speech, Language Process.*, vol. 33, pp. 3653–3665, 2025.
- [13] Alexandre Défossez, Jade Copet, Gabriel Synnaeve, and Yossi Adi, “High fidelity neural audio compression,” *Trans. Mach. Learn. Res.*, 2023.
- [14] Jade Copet, Felix Kreuk, Itai Gat, Tal Remez, David Kant, Gabriel Synnaeve, Yossi Adi, and Alexandre Défossez, “Simple and controllable music generation,” in *Proc. NeurIPS*, 2023.
- [15] Peter Knees, Ángel Faraldo, Perfecto Herrera, Richard Vogl, Sebastian Böck, et al., “Two data sets for tempo estimation and key detection in electronic dance music annotated from user corrections,” in *Proc. ISMIR*, 2015.
- [16] Jesse Engel, Cinjon Resnick, Adam Roberts, Sander Dieleman, Mohammad Norouzi, et al., “Neural audio synthesis of musical notes with WaveNet autoencoders,” in *Proc. ICML*, 2017.